

Wine Quality Classification

About dataset

The two datasets are related to red and white variants of the Portuguese "Vinho Verde" wine. The reference [Cortez et al., 2009]. Due to privacy and logistic issues, only physicochemical (inputs) and sensory (the output) variables are available (e.g. there is no data about grape types, wine brand, wine selling price, etc.).

Input variables (based on physicochemical tests):

- 1 - fixed acidity
- 2 - volatile acidity
- 3 - citric acid
- 4 - residual sugar
- 5 - chlorides
- 6 - free sulfur dioxide
- 7 - total sulfur dioxide
- 8 - density
- 9 - pH
- 10 - sulphates
- 11 - alcohol

Output variable (based on sensory data):

- 12 - quality (score between 0 and 10)

Imports

```
In [362]: import pandas as pd
import numpy as np
from sklearn import tree
from sklearn.utils import shuffle
from sklearn.preprocessing import LabelEncoder, StandardScaler, MinMaxScaler
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense
import seaborn as sns
from tensorflow import keras
```

Loading data and shuffling the dataset

```
In [363]: df = pd.read_csv("winequalityN.csv")
df = shuffle(df)
```

Out[363]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	6487.000000	6489.000000	6494.000000	6495.000000	6495.000000	6497.000000	6497.000000
mean	7.216579	0.339691	0.318722	5.444326	0.056042	30.525319	115.7445
std	1.296750	0.164649	0.145265	4.758125	0.035036	17.749400	56.5218
min	3.800000	0.080000	0.000000	0.600000	0.009000	1.000000	6.0000
25%	6.400000	0.230000	0.250000	1.800000	0.038000	17.000000	77.0000
50%	7.000000	0.290000	0.310000	3.000000	0.047000	29.000000	118.0000
75%	7.700000	0.400000	0.390000	8.100000	0.065000	41.000000	156.0000
max	15.900000	1.580000	1.660000	65.800000	0.611000	289.000000	440.0000

```
In [364]: df.head(10)
```

Out[364]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphate
5339	red	11.9	0.40	0.65	2.15	0.068	7.0	27.0	0.99880	3.06	0.6
3217	white	5.8	0.33	0.23	5.00	0.053	29.0	106.0	0.99458	3.13	0.5
3992	white	6.7	0.19	0.32	3.70	0.041	26.0	76.0	0.99173	2.90	0.5

Dataset info

As we can see we have unfortunately inconsistent data in some rows there are data gaps

In [365]:

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6497 entries, 5339 to 63
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   type                   6497 non-null   object
1   fixed acidity          6487 non-null   float64
2   volatile acidity       6489 non-null   float64
3   citric acid            6494 non-null   float64
4   residual sugar         6495 non-null   float64
5   chlorides              6495 non-null   float64
6   free sulfur dioxide    6497 non-null   float64
7   total sulfur dioxide   6497 non-null   float64
8   density                6497 non-null   float64
9   pH                     6488 non-null   float64
10  sulphates              6493 non-null   float64
11  alcohol                6497 non-null   float64
12  quality                6497 non-null   int64
dtypes: float64(11), int64(1), object(1)
memory usage: 710.6+ KB
```

In [366]:

```
Out[366]: type                0
fixed acidity              10
volatile acidity           8
citric acid                3
residual sugar             2
chlorides                  2
free sulfur dioxide        0
total sulfur dioxide       0
density                   0
pH                         9
sulphates                  4
alcohol                   0
quality                   0
dtype: int64
```

The first database preprocess

At the beginning, we delete the rows in which we have missing data

```
In [367]: classic_df = df.dropna(axis=0)

<class 'pandas.core.frame.DataFrame'>
Int64Index: 6463 entries, 5339 to 63
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   type                   6463 non-null   object
1   fixed acidity          6463 non-null   float64
2   volatile acidity       6463 non-null   float64
3   citric acid            6463 non-null   float64
4   residual sugar         6463 non-null   float64
5   chlorides              6463 non-null   float64
6   free sulfur dioxide    6463 non-null   float64
7   total sulfur dioxide   6463 non-null   float64
8   density                6463 non-null   float64
9   pH                     6463 non-null   float64
10  sulphates              6463 non-null   float64
11  alcohol                6463 non-null   float64
12  quality                6463 non-null   int64
dtypes: float64(11), int64(1), object(1)
memory usage: 706.9+ KB
```

```
In [368]: classic_df["quality"]
```

```
Out[368]: 6    2820
5    2128
7    1074
4     214
8     192
3       30
9        5
Name: quality, dtype: int64
```

Then we divide our data into X and y where X is input and y is output (quality)

```
In [369]: X = classic_df.drop(columns="quality")
y = classic_df['quality']
```

```
Out[369]:
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphate
5339	red	11.9	0.40	0.65	2.15	0.068	7.0	27.0	0.99880	3.06	0.6
3217	white	5.8	0.33	0.23	5.00	0.053	29.0	106.0	0.99458	3.13	0.5
3992	white	6.7	0.19	0.32	3.70	0.041	26.0	76.0	0.99173	2.90	0.5

In [370]:

```
Out[370]: 5339    6
          3217    5
          3992    7
          2215    6
           87    6
          1868    7
          1376    6
          2188    6
           20    8
          192    6
          2702    5
          3882    5
          1513    6
          6069    6
          6459    5
          Name: quality, dtype: int64
```

We replace the quality rating with versions with only two solutions - good and average

```
In [371]: bins = [0, 5.5, 10]
          labels = ["average", "good"]
          y = pd.cut(y, bins=bins, labels=labels)
```

```
Out[371]: 5339    good
          3217  average
          3992    good
          2215    good
           87    good
          1868    good
          1376    good
          2188    good
           20    good
          192    good
          2702  average
          3882  average
          1513    good
          6069    good
          6459  average
          Name: quality, dtype: category
          Categories (2, object): ['average' < 'good']
```

We replace average and good with 0 and 1 and wine types (red and white) also with 0 and 1 as well

```
In [372]: le = LabelEncoder()
          y = le.fit_transform(y)
```

```
Out[372]: array([1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0])
```

```
In [373]: X['type'] = le.fit_transform(X['type'])
```

```
Out[373]:
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
5339	0	11.9	0.40	0.65	2.15	0.068	7.0	27.0	0.99880	3.06	0.68
3217	1	5.8	0.33	0.23	5.00	0.053	29.0	106.0	0.99458	3.13	0.52
3992	1	6.7	0.19	0.32	3.70	0.041	26.0	76.0	0.99173	2.90	0.57
2215	1	8.5	0.28	0.34	13.80	0.041	32.0	161.0	0.99810	3.13	0.40
87	1	6.8	0.25	0.31	13.30	0.050	69.0	202.0	0.99720	3.22	0.48

Data standardization

```
In [374]: sc = StandardScaler()  
X = sc.fit_transform(X)
```

```
Out[374]: (6463, 12)
```

Data split into test and training set 30 - 70%

```
In [375]:
```

```
from sklearn.model_selection import train_test_split
```

Decision trees

We use two decision trees - one without constraints and the other with a maximum depth of three levels

```
In [376]: clf = tree.DecisionTreeClassifier()
```

```
In [377]: clf = clf.fit(x_train, y_train)
```

In [378]:

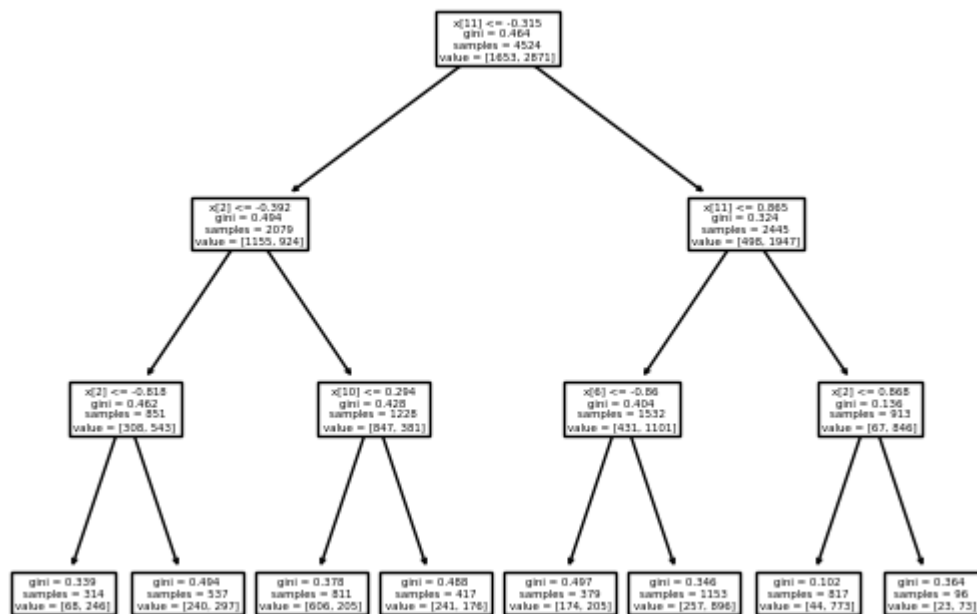
```
Out[378]: [Text(0.5338475113646644, 0.9772727272727273, 'x[11] <= -0.315\ngini = 0.464\
nsamples = 4524\nvalue = [1653, 2871]'),
Text(0.25366312533924373, 0.9318181818181818, 'x[2] <= -0.392\ngini = 0.494\
nsamples = 2079\nvalue = [1155, 924]'),
Text(0.06310569137868645, 0.8863636363636364, 'x[2] <= -0.818\ngini = 0.462\
nsamples = 851\nvalue = [308, 543]'),
Text(0.03690971593993125, 0.8409090909090909, 'x[4] <= 0.884\ngini = 0.339\n
samples = 314\nvalue = [68, 246]'),
Text(0.029491586755925458, 0.7954545454545454, 'x[6] <= -0.958\ngini = 0.41
7\nsamples = 182\nvalue = [54, 128]'),
Text(0.025511127193776007, 0.75, 'x[4] <= -0.493\ngini = 0.484\nsamples = 1
7\nvalue = [10, 7]'),
Text(0.02406368735299439, 0.7045454545454546, 'x[5] <= 2.707\ngini = 0.165\n
samples = 11\nvalue = [10, 1]'),
Text(0.022616247512212775, 0.6590909090909091, 'gini = 0.0\nsamples = 10\nva
lue = [10, 0]'),
Text(0.025511127193776007, 0.6590909090909091, 'gini = 0.0\nsamples = 1\nval
ue = [0, 1]'),
Text(0.026958567034557627, 0.7045454545454546, 'gini = 0.0\nsamples = 6\nval
```

In [379]:

```

Out[379]: [Text(0.5, 0.875, 'x[11] <= -0.315\ngini = 0.464\nsamples = 4524\nvalue = [16
53, 2871]'),
Text(0.25, 0.625, 'x[2] <= -0.392\ngini = 0.494\nsamples = 2079\nvalue = [11
55, 924]'),
Text(0.125, 0.375, 'x[2] <= -0.818\ngini = 0.462\nsamples = 851\nvalue = [30
8, 543]'),
Text(0.0625, 0.125, 'gini = 0.339\nsamples = 314\nvalue = [68, 246]'),
Text(0.1875, 0.125, 'gini = 0.494\nsamples = 537\nvalue = [240, 297]'),
Text(0.375, 0.375, 'x[10] <= 0.294\ngini = 0.428\nsamples = 1228\nvalue = [8
47, 381]'),
Text(0.3125, 0.125, 'gini = 0.378\nsamples = 811\nvalue = [606, 205]'),
Text(0.4375, 0.125, 'gini = 0.488\nsamples = 417\nvalue = [241, 176]'),
Text(0.75, 0.625, 'x[11] <= 0.865\ngini = 0.324\nsamples = 2445\nvalue = [49
8, 1947]'),
Text(0.625, 0.375, 'x[6] <= -0.86\ngini = 0.404\nsamples = 1532\nvalue = [43
1, 1101]'),
Text(0.5625, 0.125, 'gini = 0.497\nsamples = 379\nvalue = [174, 205]'),
Text(0.6875, 0.125, 'gini = 0.346\nsamples = 1153\nvalue = [257, 896]'),
Text(0.875, 0.375, 'x[2] <= 0.868\ngini = 0.136\nsamples = 913\nvalue = [67,
846]'),
Text(0.8125, 0.125, 'gini = 0.102\nsamples = 817\nvalue = [44, 773]'),
Text(0.9375, 0.125, 'gini = 0.364\nsamples = 96\nvalue = [23, 73]')]

```





```
In [380]: prediction = clf.predict(x_test)
prediction_smaller = clf_smaller.predict(x_test)

print("Accuracy on test data set with bigger tree: ", accuracy_score(prediction,
y_test))
print("Accuracy on test data set with smaller tree: ", accuracy_score(prediction_smaller,
y_test))
```


Accuracy on test data set with bigger tree: 0.7467766890149562
Accuracy on test data set with smaller tree: 0.7282104177411036

As we can see, a deeper decision tree performs better, but it takes more time

```
In [381]:
```



```
In [382]: ax = sns.heatmap(confusion_matrix(y_test, prediction_smaller), annot=True, fmt
```




Naive-Bayes

```
In [383]: model = GaussianNB()  
model.fit(x_train, y_train)
```

Accuracy on test data set: 0.6730273336771532

As we can see, it has worse accuracy than the decision trees

```
In [384]: prediction = model.predict(x_test)
```




K-nearest neighbors

First try with three neighbors

```
In [385]: knn = KNeighborsClassifier(n_neighbors=3, metric='euclidean')  
knn.fit(x_train, y_train)
```

Accuracy on test data set: 0.7607013924703455

```
In [386]: prediction = knn.predict(x_test)
```

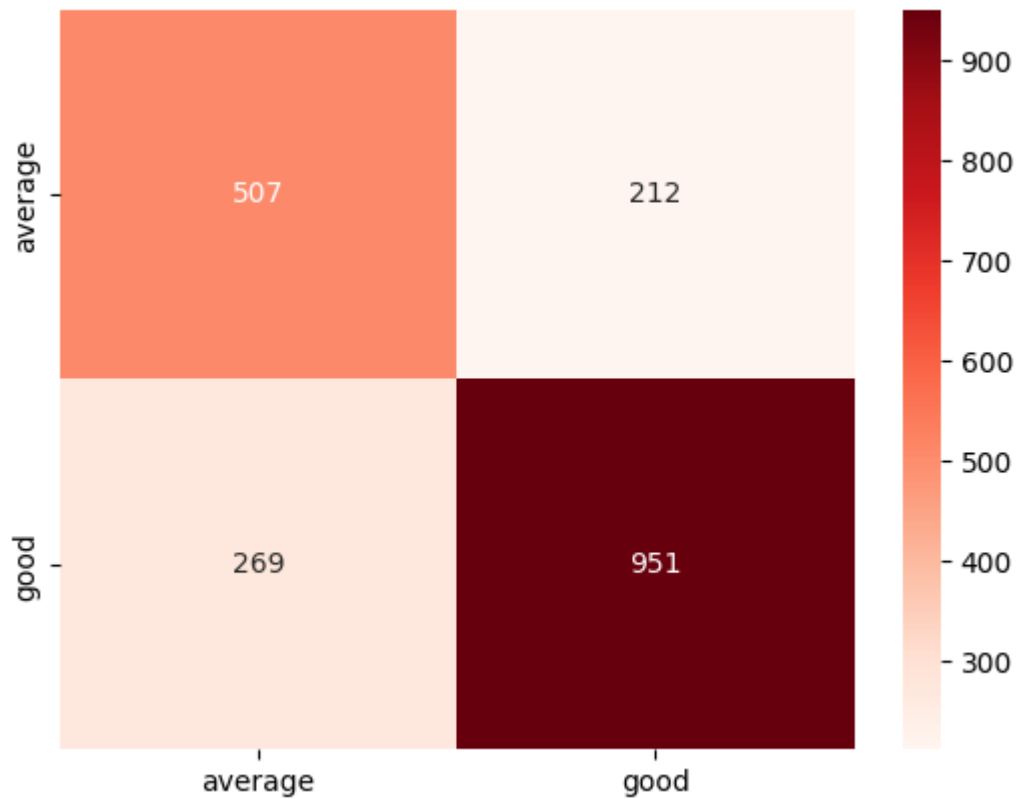


Six neighbors

```
In [387]: knn = KNeighborsClassifier(n_neighbors=6, metric='euclidean')  
knn.fit(x_train, y_train)
```

Accuracy on test data set: 0.7519339865910263

```
In [388]: prediction = knn.predict(x_test)
```




Nine neighbors

```
In [389]: knn = KNeighborsClassifier(n_neighbors=9, metric='euclidean')  
knn.fit(x_train, y_train)
```

Accuracy on test data set: 0.7488396080453842

```
In [390]: prediction = knn.predict(x_test)
```




Twelve neighbors

```
In [391]: knn = KNeighborsClassifier(n_neighbors=12, metric='euclidean')  
knn.fit(x_train, y_train)
```

Accuracy on test data set: 0.7457452294997421

```
In [392]: prediction = knn.predict(x_test)
```



As we can see, there is no big difference in the results

Neural networks


First try with MLPClassifier from sklearn with structure 6, 3, relu activation function, solver adam

```
In [393]: clf = MLPClassifier(solver='adam', alpha=1e-5, hidden_layer_sizes=(6, 3), random_state=1)
clf = clf.fit(x_train, y_train)
prediction = clf.predict(x_test)
```

Accuracy on test data set: 0.7498710675605983

```
c:\Users\Piotr Damrych\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\neural_network\_multilayer_perceptron.py:686: ConvergenceWarning:
Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
warnings.warn(
```

In [394]:



Preparation of the dataset for neural networks with keras

```
In [395]: y_train = keras.utils.to_categorical(y_train, num_classes=2)
y_test_to_validation = keras.utils.to_categorical(y_test, num_classes=2)
```

```
Out[395]: array([[0., 1.],
 [0., 1.],
 [0., 1.],
 [0., 1.],
 [1., 0.]], dtype=float32)
```

First try with 3, 6, 2 structure, relu and sigmoid activation functions, solver adam

```
In [396]: model = Sequential()
model.add(Dense(3, activation='relu', input_dim=(x_train.shape[1])))
model.add(Dense(6, activation='relu'))
model.add(Dense(2, activation='softmax'))


model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```




```
In [397]: pred = model.predict(x_test)
pred = np.argmax(pred, axis=-1)
```

```
61/61 [=====] - 0s 1ms/step
Accuracy: 0.7529654461062403
```

```
In [398]:
```



```
In [399]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Test'], loc='upper left')
```



Second try with 3, 6, 2 structure, elu and sigmoid activation functions, solver adam

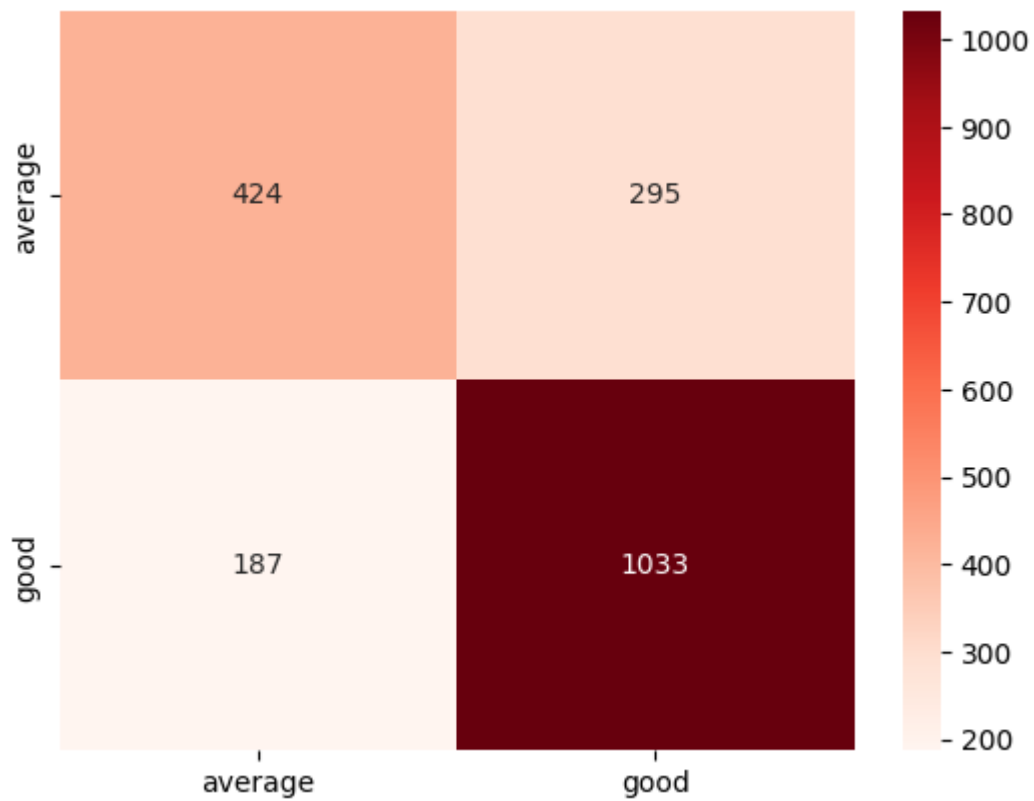

```
In [400]: model = Sequential()
model.add(Dense(3, activation='elu', input_dim=(x_train.shape[1])))
model.add(Dense(6, activation='elu'))
model.add(Dense(2, activation='sigmoid'))
```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```


```
In [401]: pred = model.predict(x_test)
pred = np.argmax(pred, axis=-1)
```

```
61/61 [=====] - 0s 1ms/step
Accuracy: 0.7514182568334193
```

In [402]:



```
In [403]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Test'], loc='upper left')
```



Third try with 3, 4, 2 structure, tanh and sigmoid activation functions, solver adam


```
In [404]: model = Sequential()
model.add(Dense(3, activation='tanh', input_dim=(x_train.shape[1])))
model.add(Dense(4, activation='tanh'))
model.add(Dense(2, activation='sigmoid'))
```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```


```
In [405]: pred = model.predict(x_test)
pred = np.argmax(pred, axis=-1)
```

```
61/61 [=====] - 0s 1ms/step
Accuracy: 0.7545126353790613
```

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100



```
In [407]: plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.title('Model loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend(['Train', 'Test'], loc='upper left')
```



Another dataset processing

This time, instead of deleting rows with empty values, we extract the mean for them

```
In [408]: different_df = df
different_df['fixed acidity'] = different_df['fixed acidity'].fillna(different_df['fixed acidity'].mean())
different_df['volatile acidity'] = different_df['volatile acidity'].fillna(different_df['volatile acidity'].mean())
different_df['citric acid'] = different_df['citric acid'].fillna(different_df['citric acid'].mean())
different_df['residual sugar'] = different_df['residual sugar'].fillna(different_df['residual sugar'].mean())
different_df['chlorides'] = different_df['chlorides'].fillna(different_df['chlorides'].mean())
different_df['pH'] = different_df['pH'].fillna(different_df['pH'].mean())
different_df['sulphates'] = different_df['sulphates'].fillna(different_df['sulphates'].mean())

<class 'pandas.core.frame.DataFrame'>
Int64Index: 6497 entries, 5339 to 63
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   type                                  6497 non-null   object
1   fixed acidity                         6497 non-null   float64
2   volatile acidity                      6497 non-null   float64
3   citric acid                          6497 non-null   float64
4   residual sugar                       6497 non-null   float64
5   chlorides                           6497 non-null   float64
6   free sulfur dioxide                  6497 non-null   float64
7   total sulfur dioxide                 6497 non-null   float64
8   density                             6497 non-null   float64
9   pH                                   6497 non-null   float64
10  sulphates                           6497 non-null   float64
11  alcohol                             6497 non-null   float64
12  quality                             6497 non-null   int64
dtypes: float64(11), int64(1), object(1)
memory usage: 710.6+ KB
```

```
In [409]: different_df['quality'] = different_df['quality'].fillna(different_df['quality'].mean())
```

```
Out[409]: 6    2836
5    2138
7    1079
4     216
8     193
3      30
9       5
Name: quality, dtype: int64
```

Preparing the dataset is the same as in the previous example with one exception

The only difference is normalization, where we use MinMaxScaler instead of StandardScaler, which converts all values to values between 0 and 1

```
In [410]: X = different_df.drop(columns="quality")
y = different_df['quality']
```

Out[410]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphate
5339	red	11.9	0.40	0.65	2.15	0.068	7.0	27.0	0.99880	3.06	0.6
3217	white	5.8	0.33	0.23	5.00	0.053	29.0	106.0	0.99458	3.13	0.5
3992	white	6.7	0.19	0.32	3.70	0.041	26.0	76.0	0.99173	2.90	0.5

```
In [411]: bins = [0, 5.5, 10]
labels = ["average", "good"]
y = pd.cut(y, bins=bins, labels=labels)
```

Out[411]:

5339	good
3217	average
3992	good
2215	good
87	good
1868	good
1376	good
2188	good
20	good
192	good
2702	average
3882	average
1513	good
6069	good
6459	average

Name: quality, dtype: category
Categories (2, object): ['average' < 'good']

```
In [412]: le = LabelEncoder()
y = le.fit_transform(y)
```

Out[412]: array([1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0])

```
In [413]: X['type'] = le.fit_transform(X['type'])
```

Out[413]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
5339	0	11.9	0.40	0.65	2.15	0.068	7.0	27.0	0.99880	3.06	0.6
3217	1	5.8	0.33	0.23	5.00	0.053	29.0	106.0	0.99458	3.13	0.5
3992	1	6.7	0.19	0.32	3.70	0.041	26.0	76.0	0.99173	2.90	0.5
2215	1	8.5	0.28	0.34	13.80	0.041	32.0	161.0	0.99810	3.13	0.4
87	1	6.8	0.25	0.31	13.30	0.050	69.0	202.0	0.99720	3.22	0.4


```
In [414]: sc = MinMaxScaler()
X = sc.fit_transform(X)
```

```
Out[414]: (6497, 12)
```

Data split into test and training set 30 - 70%

```
In [415]:
```

Decision trees

We use two decision trees - one without constraints and the other with a maximum depth of three levels

```
In [416]: clf = tree.DecisionTreeClassifier()
```

```
In [417]: clf = clf.fit(x_train, y_train)
```

```
In [418]:
```

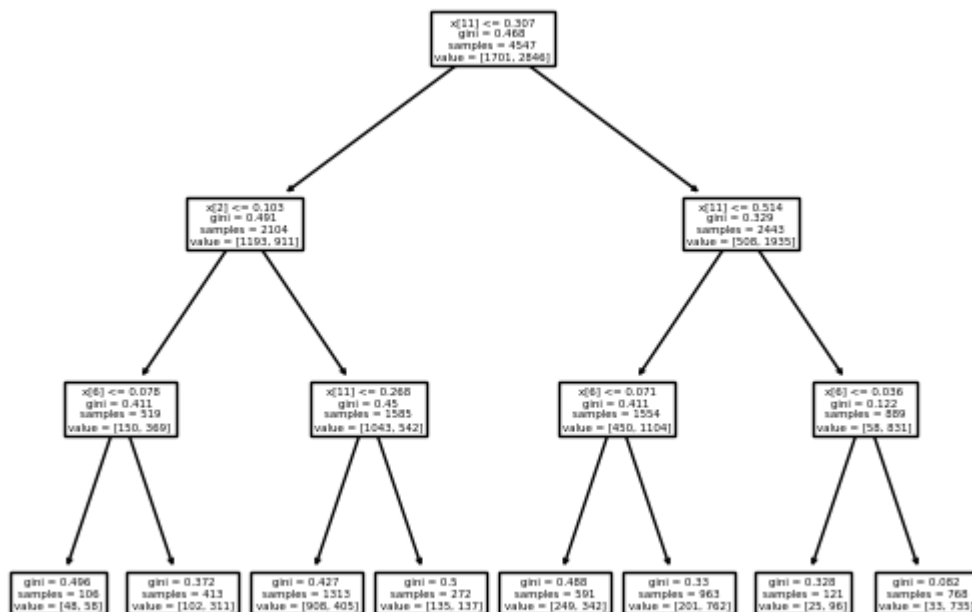
```
Out[418]: [Text(0.5299146938094853, 0.9791666666666666, 'x[11] <= 0.307\ngini = 0.468\nsamples = 4547\nvalue = [1701, 2846]'),
Text(0.21382269985330807, 0.9375, 'x[2] <= 0.103\ngini = 0.491\nsamples = 2104\nvalue = [1193, 911]'),
Text(0.04529257530955461, 0.8958333333333334, 'x[6] <= 0.078\ngini = 0.411\nsamples = 519\nvalue = [150, 369]'),
Text(0.014045462945851045, 0.8541666666666666, 'x[4] <= 0.038\ngini = 0.496\nsamples = 106\nvalue = [48, 58]'),
Text(0.0073923489188689705, 0.8125, 'x[5] <= 0.092\ngini = 0.459\nsamples = 56\nvalue = [36, 20]'),
Text(0.005913879135095177, 0.7708333333333334, 'x[8] <= 0.126\ngini = 0.426\nsamples = 52\nvalue = [36, 16]'),
Text(0.0029569395675475884, 0.7291666666666666, 'x[6] <= 0.043\ngini = 0.497\nsamples = 28\nvalue = [15, 13]'),
Text(0.0014784697837737942, 0.6875, 'gini = 0.0\nsamples = 10\nvalue = [10, 0]'),
Text(0.004435409351321382, 0.6875, 'x[5] <= 0.069\ngini = 0.401\nsamples = 18\nvalue = [5, 13]'),
Text(0.0029569395675475884, 0.6458333333333334, 'x[6] <= 0.068\ngini = 0.23\nsamples = 15\nvalue = [5, 10]')]
```

In [419]:

```

Out[419]: [Text(0.5, 0.875, 'x[11] <= 0.307\ngini = 0.468\nsamples = 4547\nvalue = [170
1, 2846]'),
Text(0.25, 0.625, 'x[2] <= 0.103\ngini = 0.491\nsamples = 2104\nvalue = [119
3, 911]'),
Text(0.125, 0.375, 'x[6] <= 0.078\ngini = 0.411\nsamples = 519\nvalue = [15
0, 369]'),
Text(0.0625, 0.125, 'gini = 0.496\nsamples = 106\nvalue = [48, 58]'),
Text(0.1875, 0.125, 'gini = 0.372\nsamples = 413\nvalue = [102, 311]'),
Text(0.375, 0.375, 'x[11] <= 0.268\ngini = 0.45\nsamples = 1585\nvalue = [10
43, 542]'),
Text(0.3125, 0.125, 'gini = 0.427\nsamples = 1313\nvalue = [908, 405]'),
Text(0.4375, 0.125, 'gini = 0.5\nsamples = 272\nvalue = [135, 137]'),
Text(0.75, 0.625, 'x[11] <= 0.514\ngini = 0.329\nsamples = 2443\nvalue = [50
8, 1935]'),
Text(0.625, 0.375, 'x[6] <= 0.071\ngini = 0.411\nsamples = 1554\nvalue = [45
0, 1104]'),
Text(0.5625, 0.125, 'gini = 0.488\nsamples = 591\nvalue = [249, 342]'),
Text(0.6875, 0.125, 'gini = 0.33\nsamples = 963\nvalue = [201, 762]'),
Text(0.875, 0.375, 'x[6] <= 0.036\ngini = 0.122\nsamples = 889\nvalue = [58,
831]'),
Text(0.8125, 0.125, 'gini = 0.328\nsamples = 121\nvalue = [25, 96]'),
Text(0.9375, 0.125, 'gini = 0.082\nsamples = 768\nvalue = [33, 735]')]

```




```
In [420]: prediction = clf.predict(x_test)
prediction_smaller = clf_smaller.predict(x_test)

print("Accuracy on test data set with bigger tree: ", accuracy_score(prediction,
y_test))
print("Accuracy on test data set with smaller tree: ", accuracy_score(prediction_smaller,
y_test))
```


Accuracy on test data set with bigger tree: 0.7764102564102564
Accuracy on test data set with smaller tree: 0.7343589743589743

As we can see, a deeper decision tree performs better, but it takes more time

In [421]:



```
In [422]: ax = sns.heatmap(confusion_matrix(y_test, prediction_smaller), annot=True, fmt
```




Naive-Bayes

```
In [423]: model = GaussianNB()  
model.fit(x_train, y_train)
```

Accuracy on test data set: 0.6805128205128205

```
In [424]: prediction = model.predict(x_test)
```




K-nearest neighbors

First try with three neighbors

```
In [425]: knn = KNeighborsClassifier(n_neighbors=3, metric='euclidean')  
knn.fit(x_train, y_train)
```

Accuracy on test data set: 0.7487179487179487

```
In [426]: prediction = knn.predict(x_test)
```

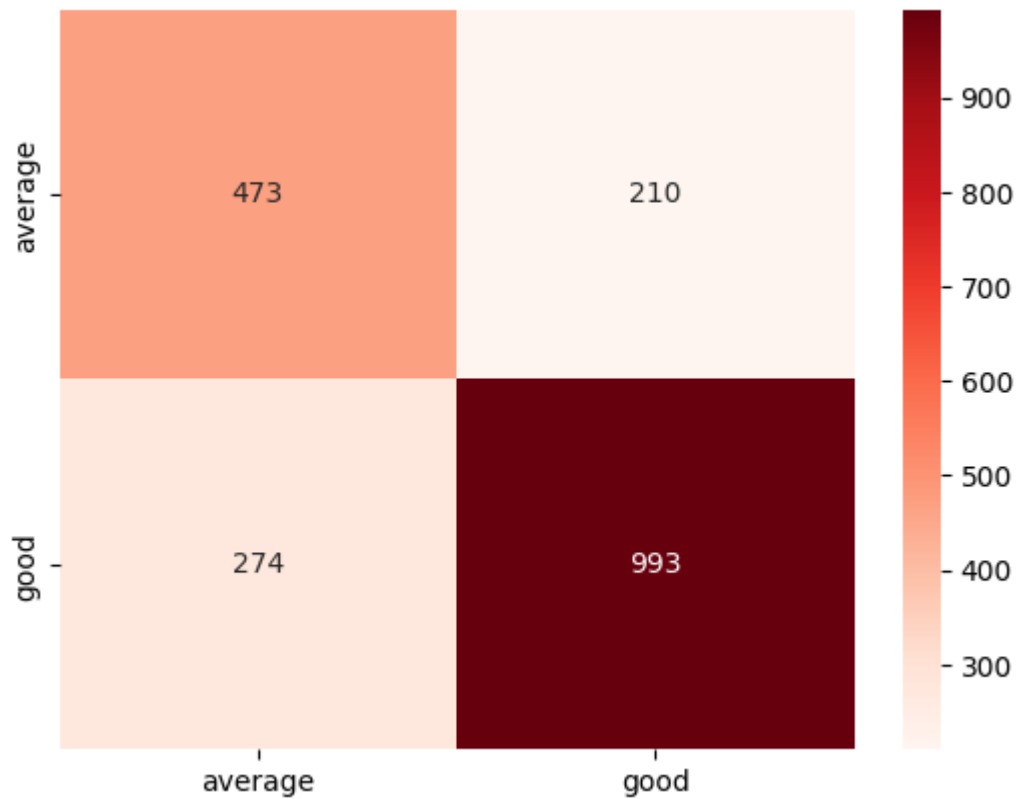


Six neighbors

```
In [427]: knn = KNeighborsClassifier(n_neighbors=6, metric='euclidean')  
knn.fit(x_train, y_train)
```

Accuracy on test data set: 0.7517948717948718

```
In [428]: prediction = knn.predict(x_test)
```




Nine neighbors

```
In [429]: knn = KNeighborsClassifier(n_neighbors=9, metric='euclidean')  
knn.fit(x_train, y_train)
```

Accuracy on test data set: 0.7538461538461538

```
In [430]: prediction = knn.predict(x_test)
```




Twelve neighbors

```
In [431]: knn = KNeighborsClassifier(n_neighbors=12, metric='euclidean')  
knn.fit(x_train, y_train)
```

Accuracy on test data set: 0.7497435897435898


```
In [432]: prediction = knn.predict(x_test)
```



As we can see, there is no big difference in the results


Neural networks

First try with MLPClassifier from sklearn with structure 6, 3, relu activation function, solver adam

```
In [433]: clf = MLPClassifier(solver='adam', alpha=1e-5, hidden_layer_sizes=(6, 3), random_state=1)
clf = clf.fit(x_train, y_train)
prediction = clf.predict(x_test)
```

Accuracy on test data set: 0.7441025641025641

In [434]:



Preparation of the dataset for neural networks with keras

```
In [435]: y_train = keras.utils.to_categorical(y_train, num_classes=2)
y_test_to_validation = keras.utils.to_categorical(y_test, num_classes=2)
```

```
Out[435]: array([[0., 1.],
[1., 0.],
[0., 1.],
[0., 1.],
[0., 1.]], dtype=float32)
```

First try with 3, 4, 2 structure, tanh and sigmoid activation functions, solver adam


```
In [436]: model = Sequential()
model.add(Dense(3, activation='tanh', input_dim=(x_train.shape[1])))
model.add(Dense(4, activation='tanh'))
model.add(Dense(2, activation='sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```


```
In [437]: pred = model.predict(x_test)
pred = np.argmax(pred, axis=-1)
```

```
61/61 [=====] - 0s 1ms/step
Accuracy: 0.7384615384615385
```

```
In [438]:
```



```
In [439]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Test'], loc='upper left')
```



Second try with 3, 6, 2 structure, elu and sigmoid activation functions, solver adam

```
In [440]: model = Sequential()
model.add(Dense(3, activation='elu', input_dim=(x_train.shape[1])))
model.add(Dense(6, activation='elu'))
model.add(Dense(2, activation='sigmoid'))
```


```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
In [441]: pred = model.predict(x_test)
pred = np.argmax(pred, axis=-1)
```

```
61/61 [=====] - 0s 1ms/step
Accuracy: 0.7492307692307693
```



```
In [443]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Test'], loc='upper left')
```



Third try with 3, 6, 2 structure, relu and sigmoid activation functions, solver adam



```
In [444]: model = Sequential()
model.add(Dense(3, activation='relu', input_dim=(x_train.shape[1])))
model.add(Dense(6, activation='relu'))
model.add(Dense(2, activation='sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```


```
In [445]: pred = model.predict(x_test)
pred = np.argmax(pred, axis=-1)
```

```
61/61 [=====] - 0s 2ms/step
Accuracy: 0.7502564102564102
```

In [446]:



```
In [447]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Train', 'Test'], loc='upper left')
```



Summary

All tested classifiers obtained similar final results on both processing of the adataset. The results were around 75%, but the highest percentage was obtained by the deep decision tree, which with the second version of the dataset was about 77%. The worst was the naive bayes classifier, which in both cases obtained < 70%

Bibliography

<https://www.kaggle.com/datasets/rajyellow46/wine-quality> (<https://www.kaggle.com/datasets/rajyellow46/wine-quality>)

